

# OS Project 2

## Scheduling in Linux

Advisor: Prof. Tei-Wei Kuo

TAs: Han-Yi Lin, Yu-Chen Lin,  
Yi-Shen Chen, Yu-Chuan Chang

# Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
  - ▶ Part I: Invoke FIFO Scheduler
  - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

# Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
  - ▶ Part I: Invoke FIFO Scheduler
  - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

# Process Life Cycle

- ▶ A process is **not** always ready to run.
- ▶ The scheduler must know the status of every process in the system when switching between tasks.
- ▶ A process may have one of the following **states**:
  - ▶ **Running** — The process is **executing at the moment**.
  - ▶ **Ready** — The process is able to run but is not allowed to **because the CPU is allocated to another process**. The scheduler can select the process at the next task switch.
  - ▶ **Waiting** — The process is sleeping and cannot run because it is **waiting for an external event**. The scheduler cannot select the process at the next task switch.
- ▶ The system saves all processes in a process table.

# The Need of the Scheduler

- ▶ A unique description of each process is held in memory and is linked with other processes by means of several structures.
- ▶ This is the situation facing the scheduler, whose task is **to share CPU time between the programs to create the illusion of concurrent execution.**
- ▶ This task is split into two different parts —
  - ▶ One relating to the **scheduling policy** and
  - ▶ The other to context switching

# Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
  - ▶ Part I: Invoke FIFO Scheduler
  - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

# Scheduling in Linux (1/2)

- ▶ The schedule function is the starting point to an understanding of scheduling operations.
- ▶ It is defined in “`kernel/sched.c`” and is one of the most frequently invoked functions in the kernel code.
- ▶ Not only priority scheduling but also two other soft real-time policies required by the POSIX standard are implemented.
  - ▶ E.g., completely fair scheduling, real-time scheduling and scheduling of the idle task, etc.

# Scheduling in Linux (2/2)

- ▶ The scheduler uses a series of data structures to sort and manage the processes in the system.
- ▶ Scheduling can be activated in two ways:
  - ▶ **Main scheduler**: Either directly if a process goes to sleep or wants to yield the CPU for other reasons,
  - ▶ **Periodic scheduler**: Or by a periodic mechanism that is run with constant frequency to check from time to time if switching tasks is necessary
- ▶ Generic scheduler = Main + Periodic schedulers



Generic Scheduler

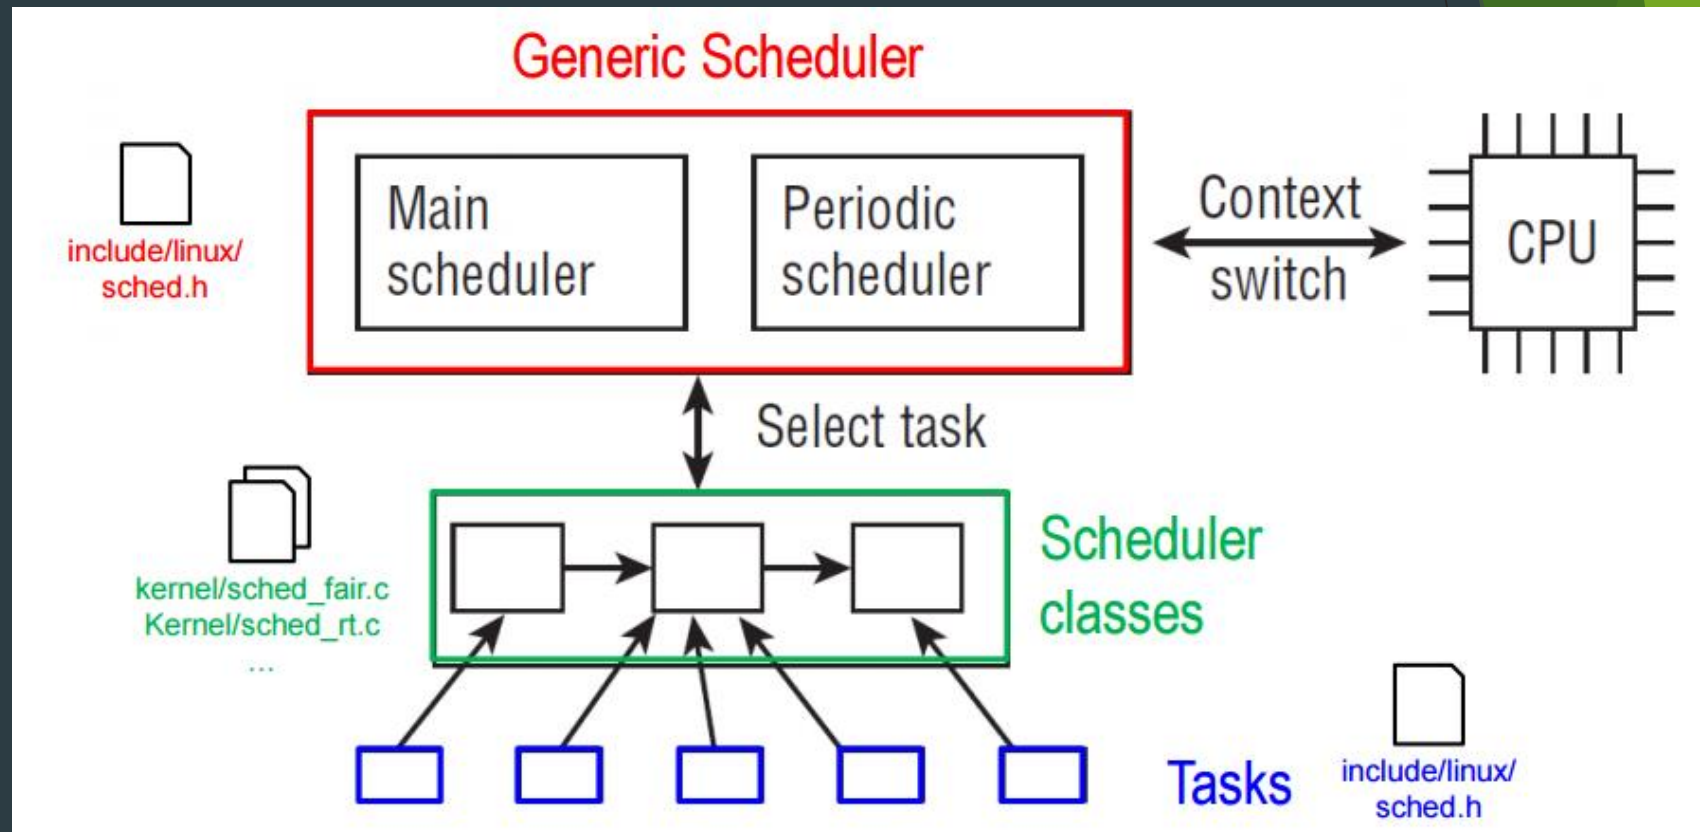
Scheduler Classes

Task

Task

Task

# Overview of the Scheduling Subsystem in Linux



# Task Representation

- In Linux, all concerned with processes and programs are built around a data structure: **task\_struct**.

<sched.h>

```
struct task_struct {  
    volatile long state;      /* -1 unrunnable, 0 runnable, >0 stopped */  
    void *stack;  
    atomic_t usage;  
    unsigned long flags;      /* per process flags, defined below */  
    unsigned long ptrace;  
    int lock_depth;           /* BKL lock depth */  
  
    int prio, static_prio, normal_prio;  
    struct list_head run_list;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
  
    ... see more in "include/linux/sched.h"
```

# Policy Designation

```
<sched.h>
struct task_struct {
    ...
    int prio, static_prio, normal_prio;
    unsigned int rt_priority;

    struct list_head run_list;
    const struct sched_class *sched_class;
    struct sched_entity se;

    unsigned int policy;
    cpumask_t cpus_allowed;
    unsigned int time_slice;

    ...
}
```

# Scheduler Classes (1 / 3)

- ▶ Scheduler classes provide the connection between the **generic scheduler** and **individual scheduling methods**.
  - ▶ They are represented by several **function pointers** collected in a special data structure.
  - ▶ Each operation that can be requested by the global scheduler is represented by one pointer.
- ▶ This allows for creation of the generic scheduler **without** any knowledge about the internal working of different scheduler classes.

# Scheduler Classes (2/3)

- An instance of **struct sched\_class** must be provided for each scheduling class.


```
<sched.h>
```

```
struct sched_class {  
    const struct sched_class *next;  
  
    void (*enqueue_task) (struct rq *rq, struct task_struct *p, int wakeup);  
    void (*dequeue_task) (struct rq *rq, struct task_struct *p, int sleep);  
    void (*yield_task) (struct rq *rq);  
  
    void (*check_preempt_curr) (struct rq *rq, struct task_struct *p);  
  
    struct task_struct * (*pick_next_task) (struct rq *rq);  
    void (*put_prev_task) (struct rq *rq, struct task_struct *p);  
    void (*set_curr_task) (struct rq *rq);  
    void (*task_tick) (struct rq *rq, struct task_struct *p);  
    void (*task_new) (struct rq *rq, struct task_struct *p);  
};
```

# Scheduler Classes (3/3)

- ▶ **enqueue\_task**: adds a new process to the run queue. This happens when a process changes from a sleeping into a runnable state.
- ▶ **dequeue\_task**: provides the inverse operation: It takes a process off a run queue. Naturally, this happens when a process switches from a runnable into an un-runnable state, or when the kernel decides to take it off the run queue for other reasons.
- ▶ **yield\_task** : when a process wants to **relinquish control of the processor voluntarily**, it can use the sched\_yield system call. This triggers yield\_task to be called in the kernel.
- ▶ **pick\_next\_task**: selects the next task that is supposed to run
- ▶ **task\_tick**: is called by the **periodic scheduler** each time it is activated.

# Relationships between Generics Functions and Process States





# Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
  - ▶ Part I: Invoke FIFO Scheduler
  - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References



# Linux Scheduling Policies

- ▶ Linux Scheduling Policies
  - ▶ Normal Scheduling policies (Non-real-time)
    - ▶ SCHED\_OTHER, SCHED\_BATCH, SCHED\_IDLE.
  - ▶ Real-Time policies
    - ▶ SCHED\_FIFO, SCHED\_RR.
- ▶ The default scheduling policy is **non-real-time**.
- ▶ In this part, using Linux **real-time scheduling policy (FIFO)** to schedule threads in a process.

# Part I: Invoke FIFO Scheduler

- ▶ Write a C program (sched\_test.c) to create two threads.
- ▶ Each thread will print who is running and busy for 0.5 second.
- ▶ Run the program by default time-sharing schedule policy and show the result.

Ex. `$ ./sched_test`

- ▶ Run the program by real-time scheduling policy (FIFO) and show the result.

Ex. `$ ./sched_test SCHED_FIFO`

```
1  int main() {
2      set CPU affinity//all threads run on the same core
3      invoke FIFO_SCHED
4
5      for(i=0;i<2;i++)
6          thread_create(i)
7          print "Thread i was created"
8      for(i=0;i<2;i++)
9          thread_join(i)
10 }
11 thread_func() {
12     for(i=0;i<3;i++)
13         print "Thread # is running"
14         busy 0.5 second
15 }
```

# Result

```
OS@VM:~$ sudo ./sched_test
Thread 1 was created
Thread 2 was created
Thread 2 is running
Thread 1 is running
Thread 2 is running
Thread 1 is running
Thread 1 is running
Thread 2 is running
OS@VM:~$ sudo ./sched_test SCHED_FIFO
Thread 1 was created
Thread 2 was created
Thread 1 is running
Thread 1 is running
Thread 1 is running
Thread 2 is running
Thread 2 is running
Thread 2 is running
```

# Hint for Part I

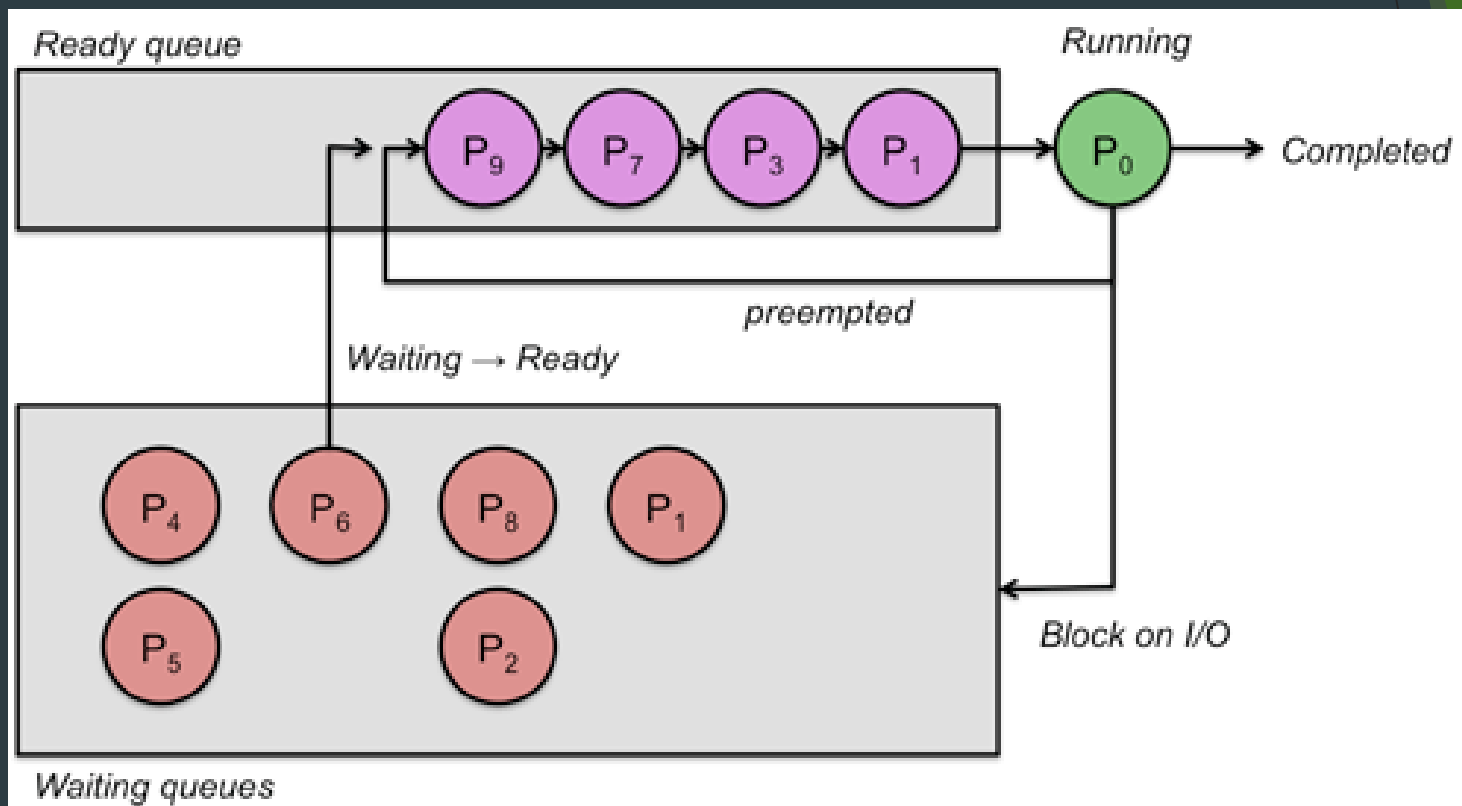
- ▶ Set CPU affinity
- ▶ `sched_setscheduler();`
- ▶ The policy corresponding value define in `/include/linux/sched.h`
- ▶ Set the priority of real-time process (`sched_param *param`)
- ▶ The permission to run real-time process

# Outline

- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
  - ▶ Part I: Invoke FIFO Scheduler
  - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

# Part II: Weighted Round Robin Scheduling (1/2)

- Processes are dispatched in a FIFO sequence but each process is allowed to run for only a limited amount of time, a.k.a., **time-slice** or **quantum**.



# Part II: Weighted Round Robin Scheduling (2/2)

- ▶ Implement `kernel/sched_weighted_rr.c`
  - ▶ `enqueue_task_weighted_rr()`
  - ▶ `dequeue_task_weighted_rr()`
  - ▶ `yield_task_weighted_rr()`
  - ▶ `pick_next_task_weighted_rr()`
  - ▶ `task_tick_weighted_rr()`
- ▶ Tasks with higher weights can finish their jobs earlier by having larger time slices.

How to add a custom scheduler  
into Linux?



# Generic Scheduler Side (1/3)

In “include/linux/sched.h”,

- Add **#define SCHED\_WEIGHTED\_RR 6** - to define your weighted rr policy

```
32  /*
33   * Scheduling policies
34   */
35  #define SCHED_NORMAL          0
36  #define SCHED_FIFO            1
37  #define SCHED_RR              2
38  #define SCHED_BATCH           3
39  /* SCHED_ISO: reserved but not implemented yet */
40  #define SCHED_IDLE            5
41  //+ OS Proj2: weighted_rr
42  #define SCHED_WEIGHTED_RR     6
```

# Generic Scheduler Side (2/3)

In “kernel/sched.c”

- Modify `__setscheduler()`, and `__sched_setscheduler()` functions - to let the generic scheduler can recognize your weighted rr scheduler

```
6515      //+ OS Proj2: weighted_rr
6516      case SCHED_WEIGHTED_RR:
6517          p->sched_class = &weighted_rr_sched_class;
6518          break;
6519      }
```

# Generic Scheduler Side (3/3)

In “struct rq” of “kernel/sched.c”

- Add **struct weighted\_rr\_rq weighted\_rr** - to specify the run queue for your weighted rr

```
506 struct rq {  
507     ...  
508     struct cfs_rq cfs;  
509     //+ OS Proj2: weighted_rr  
510     struct weighted_rr_rq weighted_rr;  
511     struct rt_rq rt;  
512     ...  
}
```

Note that struct rq - the generic per-CPU run queue structure. However, this is **NOT** the queue structure you will work with. Rather, this structure contains a more specific run queue type for different scheduler classes.

# Scheduler Classes Side (1/3)

As well in “kernel/sched.c”

- ▶ Define `weighted_rr_rq` structure, which should contain
  - ▶ `struct list_head queue` - to denote the actual run queue for your weighted rr scheduler
  - ▶ `unsigned long nr_running` - to denote the number of processes which are now in the run queue

```
424 //+ OS Proj2: weighted_rr
425 struct weighted_rr_rq {
426     struct list_head queue;
427     unsigned long nr_running;
```

# Scheduler Classes Side (2/3)

In “kernel/sched.c”,

- Declare `int weighted_rr_time_slice` - to define the time slice for your weighted rr scheduling policy

```
1934 //+ OS Proj2: weighted_rr
1935 int weighted_rr_time_slice
```

```
7227 //+ OS Proj2: weighted_rr
7228 SYSCALL_DEFINE1(sched_weighted_rr_setquantum, unsigned int, quantum)
7229 {
7230     weighted_rr_time_slice = quantum;
7231     return;
7232 }
```

# Scheduler Classes Side (3/3)

In “kernel/sched\_weighted\_rr.c”

- ▶ Accomplish the implementation of weighted rr scheduler
  - ▶ Recall that an instance of **struct sched\_class** must be provided for each scheduling class.

```
243  const struct sched_class weighted_rr_sched_class = {
244      .next                = &idle_sched_class,
245      .enqueue_task        = enqueue_task_weighted_rr,
246      .dequeue_task        = dequeue_task_weighted_rr,
247      .yield_task          = yield_task_weighted_rr,
248
249      .check_preempt_curr  = check_preempt_curr_weighted_rr,
250
251      .pick_next_task       = pick_next_task_weighted_rr,
252      .put_prev_task       = put_prev_task_weighted_rr,
253      ...

```

# Task Side

In “struct task\_struct” of “include/linux/sched.h”, add

- ▶ Declare **unsigned int weighted\_rr\_task\_time\_slice** - to denote the current time slice for this task
- ▶ Declare **struct list\_head weighted\_rr\_list\_item** - to denote the list item which will be inserted into the run queue of weighted\_rr

```
1219 struct task_struct {
1220     ...
1221     //+ OS Proj2: weighted_rr
1222     unsigned int task_time_slice;
1223     unsigned int weighted_time_slice;
1224     ...
1225     //+ OS Proj2: weighted_rr
1226     struct list_head weighted_rr_list_item;
```

# “Lazy Package”

- ▶ The lazy package includes

- ▶ <http://newslab.csie.ntu.edu.tw/course/OS2018/files/project/PJ2-linux-2.6.32.60.tar.gz>

- ▶ Six modified files (don't modify, but **read it**)

- ▶ include/linux/sched.h, kernel/sched.c, kernel/sched\_fair.c,  
include/linux/syscalls.h,  
arch/x86/kernel/syscall\_table\_32.S,  
arch/x86/include/asm/unistd\_32.h

- ▶ sched\_weighted\_rr.c (**incomplete, your job!**)

- enqueue\_task\_weighted\_rr(), dequeue\_task\_weighted\_rr()

- yield\_task\_weighted\_rr(), pick\_next\_task\_weighted\_rr()

- task\_tick\_weighted\_rr()



# Testing Program

In linux-2.6.32.60\test\_weighted\_rr\test\_weighted\_rr.c

- ▶ The test program will first allocate a write buffer with size **b**.
- ▶ Then, the test program will create **n** user threads, each of which will write a unique character (e.g., a) into the buffer over and over.
  - ▶ Note that, every threads will write the same number of characters in to the buffer, based on the buffer size.
- ▶ Moreover, you can assign the scheduling policy, and the weighted\_rr\_time\_slice **t**.

Note that, when dumping the write buffer, the test program will aggregate the consecutive characters into one symbol.

# Possible Results

```
▶ ./test_weighted_rr weighted_rr t 5 5000000
```

```
> ./test_weighted_rr weighted_rr 10 5 500000000  
sched_policy: 6, quantum: 10, num_threads: 5, buffer_size: 500000000  
abcdeabcdeabcdeabcdabcbcabcbcababababababababa
```

'e' finish    'd' finish    'c' finish

# test\_weighted\_rr.c

```
for (i = 0; i < num_threads; i++)
{
    ...
    syscall (SYS_weighted_rr_setquantum, quantum);
    pthread_create(&threads[i], &attr, run, (void *)targs);
    quantum*=2;
}
```

# Scoring of Project 2

- ▶ Part I: Implementation of a program to invoke FIFO scheduler (30%)
- ▶ Part II: Implementation of the below **FIVE** incomplete functions in “sched\_weighted\_rr.c” (40%)
  - enqueue\_task\_weighted\_rr(), dequeue\_task\_weighted\_rr()
  - yield\_task\_weighted\_rr(), pick\_next\_task\_weighted\_rr()
  - task\_tick\_weighted\_rr()
- ▶ Report (30%)
  - ▶ Your implementation details and results
  - ▶ At most 4 pages

# Outline

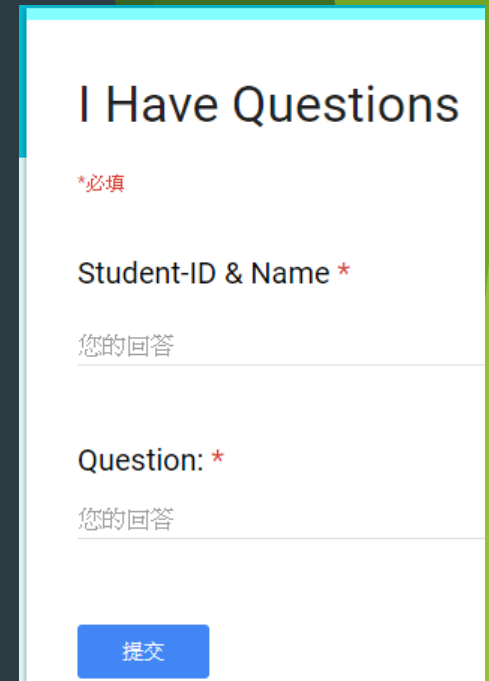
- ▶ Review: Life of Process
- ▶ Scheduling in Linux
- ▶ Implementation
  - ▶ Part I: Invoke FIFO Scheduler
  - ▶ Part II: Weighted Round Robin Scheduler
- ▶ Submission Rules
- ▶ References

# Submission Rules

- ▶ Project deadline: 2018/05/11 23:59
- ▶ Upload to FTP Server
  - ▶ IP: 140.112.28.143
  - ▶ Port: 21
  - ▶ Account: os2018
  - ▶ Password: ktw2018os
- ▶ Be packed as one file named “OSPJ2\_Team##\_v##.zip ”
  - ▶ Report.pdf
  - ▶ Part1(directory)
    - +--- sched\_test.c
  - ▶ Part2(directory)
    - +--- sched\_weighted\_rr.c
- ▶ **DO NOT COPY THE HOMEWORK**

# Contact TAs

- ▶ If you have any question about the project, please feel free to contact TAs.
- ▶ I have questions:  
<https://goo.gl/forms/39eB4ex4w3EX7l4K2>
- ▶ Video:  
<http://newslab.csie.ntu.edu.tw/course/OS2018/PJ2.html>
- ▶ Han-Yi Lin: [d03922006@csie.ntu.edu.tw](mailto:d03922006@csie.ntu.edu.tw)  
Yu-Chen Lin: [f04922077@csie.ntu.edu.tw](mailto:f04922077@csie.ntu.edu.tw)  
Yi-Shen Chen: [d05922009@csie.ntu.edu.tw](mailto:d05922009@csie.ntu.edu.tw)  
Yu-Chuan Chang: [r05922057@csie.ntu.edu.tw](mailto:r05922057@csie.ntu.edu.tw)



The image shows a screenshot of a Google Form titled "I Have Questions". The form is in Chinese and includes the following elements:

- Title: I Have Questions
- Required field indicator: \*必填
- Field 1: Student-ID & Name \*
- Field 1 input: 您的回答
- Field 2: Question: \*
- Field 2 input: 您的回答
- Submit button: 提交

# References

- ▶ Reference Book

- ▶ Professional Linux® Kernel Architecture, Wolfgang  
Mauerer, Wiley Publishing, Inc.

- ▶ Process Scheduling

- ▶ <https://www.cs.rutgers.edu/~pxk/416/notes/07-scheduling.html>