

OS Project 2

Hints

Advisor: Prof. Tei-Wei Kuo

Speaker: Yu-Chen Lin

Part I

Hints 1/3

- Set CPU affinity

Hints:


- Use the function `sched_setaffinity()` to specify one core that can be used by the scheduler and the two threads
- There are several parameters should be initialized before the `sched_setaffinity()` function call

```
thread function:
    for i=1 to 3
        print "Thread # is running"
        busy 0.5 second

main function:
    setp 1: set CPU affinity
    step 2: invoke FIFO scheduler

    for i=1 to 2
        pthread_create(i)
        print "Thread # was created"

    for i=1 to 2
        pthread_join(i)
```



Hints 2/3

- Invoke FIFO scheduler

Hints:

- Use the function **`sched_setscheduler()`** to change the scheduling policy
- There are several parameters should be initialized before the **`sched_setscheduler()`** function call

```
thread function:
    for i=1 to 3
        print "Thread # is running"
        busy 0.5 second


main function:

    setp 1: set CPU affinity

    step 2: invoke FIFO scheduler

    for i=1 to 2
        pthread_create(i)
        print "Thread # was created"

    for i=1 to 2
        pthread_join(i)
```




Hints 3/3

- Busy waiting

Discussion:

- `sleep(500)???`



```
thread function:
  for i=1 to 3
    print "Thread # is running"
    busy 0.5 second

main function:

  setp 1: set CPU affinity

  step 2: invoke FIFO scheduler

  for i=1 to 2
    pthread_create(i)
    print "Thread # was created"

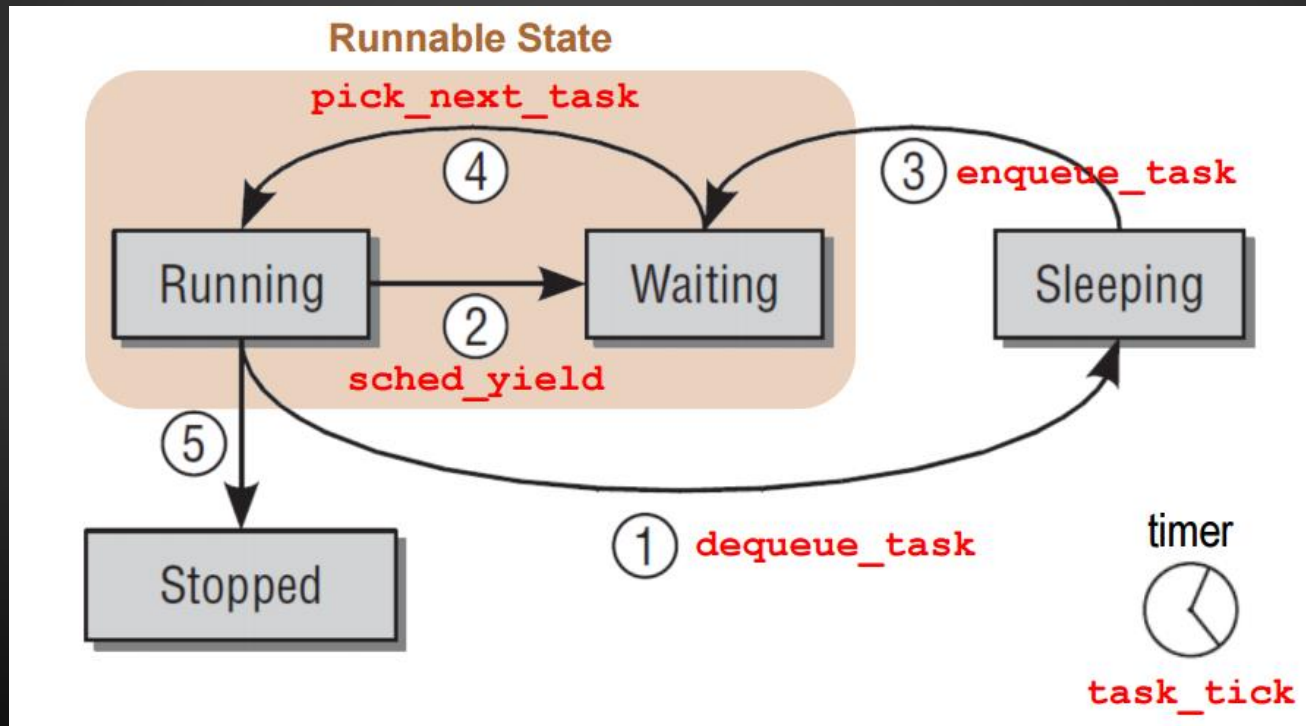
  for i=1 to 2
    pthread_join(i)
```

Part II

In linux-2.6.32.60/kernel/sched_weighted_rr.c

- Accomplish the five functions of weighted rr scheduler
 - static void **enqueue_task_weighted_rr**static void
(struct rq *rq, struct task_struct *p, int wakeup, bool b)
 - static void **dequeue_task_weighted_rr**
(struct rq *rq, struct task_struct *p, int sleep)
 - static void **yield_task_weighted_rr** (struct rq *rq)
 - static struct task_struct ***pick_next_task_weighted_rr** (struct rq *rq)
 - static **void task_tick_weighted_rr**
(struct rq *rq, struct task_struct *p, int queued)

Relationships between Generics Functions and Process States



Hints 1/4

- static void **enqueue_task_weighted_rr**
(struct rq *rq, struct task_struct *p, int wakeup, bool b)
- static void **dequeue_task_weighted_rr**
(struct rq *rq, struct task_struct *p, int sleep)

Hints:

- Use functions **list_add_tail()** and **list_del()** to enqueue and dequeue task_struct *p
- Remember to update the **rq->weighted_rr.nr_running** value after enqueueing/dequeueing

Hints 2/4

- static void **yield_task_weighted_rr** (struct rq *rq)

Hint:

- Use the function **list_move_tail()** to put the current task (**rq->curr**) to the end of the run list

Hints 3/4

- static `void task_tick_weighted_rr`
(struct rq *rq, struct task_struct *p, int queued)

Hints:

- `task_tick` is called by the periodic scheduler each time it is activated
- First, `task_time_slice` value of the task p minus one
- Once `task_time_slice` value of the task p is zero
 - ① reset `task_time_slice` of the task p
 - ② call the function `set_tsk_need_resched(q)`
 - ③ yield/requeue the task p

Hints 3/4 (cont.)

In linux-2.6.32.60/include/linux/sched.h

- **task_time_slice:**

- record the consumption of time slice

- **weighted_time_slice:**

- how much time should be supplied when reset

task_time_slice is according to the **weighted_time_slice**

```
1219 struct task_struct {
1220     ...
1221     //+ RTS Proj2: weighted_rr
1222     unsigned int task_time_slice;
1223     unsigned int weighted_time_slice;
1224     ...
1225     //+ RTS Proj2: weighted_rr
1226     struct list_head weighted_rr_list_item;
```

Hints 3/4 (cont.)

In linux-2.6.32.60/kernel/sched.c

```
static void __sched_fork(struct task_struct *p)
{
    ...
    //+ OS Proj2: weighted_rr
    INIT_LIST_HEAD(&p->weighted_rr_list_item);
    p->task_time_slice = weighted_rr_time_slice;
    p->weighted_time_slice = weighted_rr_time_slice;
    ...
}

//+ OS Proj2: weighted_rr
SYSCALL_DEFINE1(sched_weighted_rr_setquantum, unsigned int, quantum)
{
    weighted_rr_time_slice = quantum;
    return;
}
```

Hints 4/4

- static struct task_struct ***pick_next_task_weighted_rr**
(struct rq *rq)
- **pick_next_task** selects the next task that is supposed to run, while put_prev_task is called before the currently executing task is replaced with another one

Hints:

- If **weighted_rr.queue** is empty, return NULL
(determined by **weighted_rr.nr_running** value)
- Otherwise, use the function **list_first_entry()** to obtain and return the first entry/task in **weighted_rr.queue**

Contact TAs

- If you have any question about the project, please feel free to contact TAs.

- I have questions:

<https://goo.gl/forms/39eB4ex4w3EX7I4K2>

- Video:

http://newslab.csie.ntu.edu.tw/course/OS2018/PJ2_Hint.html

- Han-Yi Lin: d03922006@csie.ntu.edu.tw
Yu-Chen Lin: f04922077@csie.ntu.edu.tw
Yi-Shen Chen: d05922009@csie.ntu.edu.tw
Yu-Chuan Chang: r05922057@csie.ntu.edu.tw

I Have Questions

*必填

Student-ID & Name *

您的回答

Question: *

您的回答

提交

References

- Reference Book
 - Professional Linux® Kernel Architecture, Wolfgang Maurer, Wiley Publishing, Inc.
- Process Scheduling
 - <https://www.cs.rutgers.edu/~pxk/416/notes/07-scheduling.html>